

Processing Large Datasets

I would first set up the S3 buckets which will store the input .csv file and also the resulting output from the Glue job. Assuming this is a one time job and I know the schema, I wouldn't need to use a Glue crawler to catalog the data, but instead will define the schema directly in the Glue Pyspark script. Then I would start with writing the script using Glue script editor, pointing to the S3 location for the input file. After finishing the script, I would test it out with a small sample size dataset first to make sure there is no obvious errors. If no errors, then I would proceed to run the Glue job on the actual dataset, and proceed to monitor the progress through cloudwatch.